

data-visualization

June 5, 2018

```
In [39]: import matplotlib.pyplot as plt
         from scipy import stats
         import numpy as np
         import pandas as pd
         %matplotlib inline
```

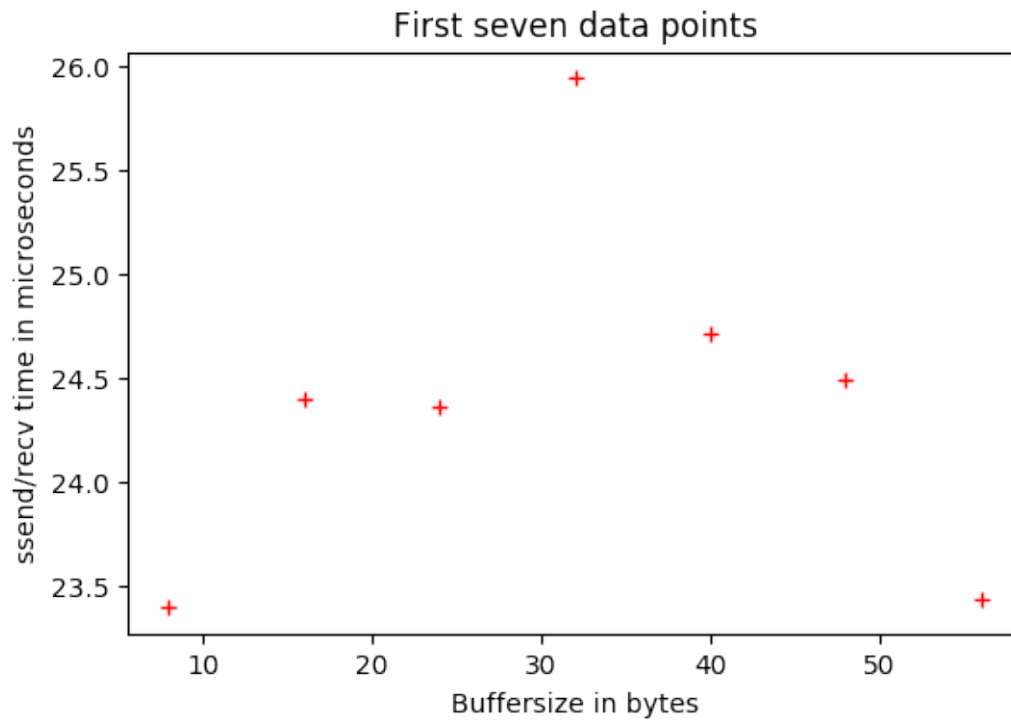
Times are measured on grid and saved in csv format. First columns is the buffersize in bytes followed by the average of send/recv time. Send messages specifies how many times the send/recv roundtrips have been measured.

```
In [81]: !head data-t40-n1000-grid.csv
```


```
buffersize;t_average;send messages
8;3.68643e-06;1000
16;4.1635e-06;1000
24;2.64752e-06;1000
40;4.71246e-06;1000
64;3.13354e-06;1000
104;4.00949e-06;1000
168;3.96705e-06;1000
272;4.26745e-06;1000
440;4.34601e-06;1000
```

```
In [25]: df = pd.read_csv('data-t40-n100000-grid.csv', sep=';')
```


```
In [36]: data_head = df[:7]
         plt.figure(dpi=100)
         plt.title('First seven data points')
         plt.xlabel('Buffersize in bytes')
         plt.ylabel('ssend/recv time in microseconds')
         plt.plot(data_head.buffersize, [x * 10e6 for x in data_head.t_average], 'r+');
```



```
In [31]: data_head = df[7:]  
plt.figure(dpi=100)  
plt.xlabel('Buffersize in bytes')  
plt.ylabel('ssend/rcv time in microseconds')  
plt.plot(data_head.buffersize, [x * 10e6 for x in data_head.t_average], 'r+');
```



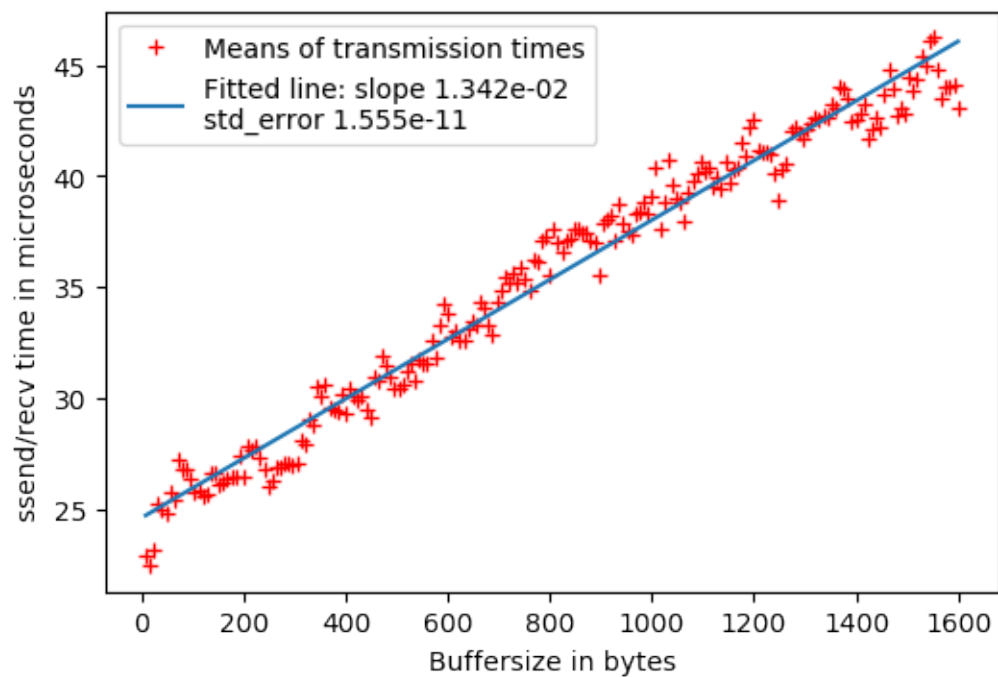
```
In [33]: df = pd.read_csv('data-lin-t200-n100000.csv', sep=';')
plt.figure(dpi=100)
plt.xlabel('Buffer size in bytes')
plt.ylabel('ssend/rcv time in microseconds')
plt.plot(df.buffer_size, [x * 10e6 for x in df.t_average], 'r+');
```



```

In [82]: df = pd.read_csv('data-lin-t200-n1000000.csv', sep=';')
        slope, intercept, r_value, p_value, std_err = stats.linregress(df.buffersize, df.t_average)
        plt.figure(dpi=100)
        plt.xlabel('Buffersize in bytes')
        plt.ylabel('ssend/rcv time in microseconds')
        plt.plot(df.buffersize, [x * 10e6 for x in df.t_average], 'r+', label="Means of transmission times")
        plt.plot(df.buffersize, (intercept+slope*df.buffersize)*10e6,
                  label="Fitted line: slope {:.3e}\nstd_error {:.3e}".format(slope*10e6, std_err))
        plt.legend();

```



```

In [83]: throughput = (df.buffersize / 1024**2) / df.t_average
        plt.figure(dpi=100)
        plt.xlabel('Buffersize in bytes')
        plt.ylabel('Throughput in MB/s')
        plt.plot(df.buffersize, throughput, 'r+');

```

